

Impact of Virtualization on Server Power Consumption

Adam Lewis Chris Thompson Nian-Feng Tzeng D’mitri Perkins

Center for Advanced Computer Studies
The University of Louisiana at Lafayette
Lafayette, Louisiana 70508

{awlewis,cct2361,tzeng,perkins}@cacs.louisiana.edu

Abstract

Large data centers host hundreds or even thousands of servers that require complex and costly cooling solutions. One approach to reducing this cost and complexity is to take advantage of virtualization techniques to run multiple instances of server software on single physical servers to reduce the overall power consumption. This paper uses experiments based upon the methodology of statistical design of experiments to investigate some of the characteristics of virtualization and to evaluate the power consumption of virtual servers.

1. Introduction

The heart of the infrastructure of today’s information economy is the data centers containing the servers that support most Internet services. Such centers host hundreds or even thousands of servers running off-the-shelf hardware. The power demands of these installations force firms to install complex and costly cooling solutions to efficiently move the heat away from servers so as to avoid reliability problems.

Providing proper cooling has become more difficult as the density of the off-the-shelf hardware has increased with tighter packaging, decreasing form factors, and increasing performance demands. The introduction of multi-core processors combined together in clusters of traditional and blade servers has greatly increased the power demands of servers.

Two approaches have been taken towards managing the power consumption of microprocessors. One approach has been the application of dynamic voltage scaling: reducing the voltage (and as a result, the frequency) of the processor at the cost of slower program execution. Other approaches (such as taken by Intel with the Pentium 4 processors) have utilized a stop-go approach: activate and deactivate the processor as demanded by workload.

These approaches have mostly been applied to battery-powered devices. They adjust components to lowest power-mode that does not over-compromise performance. These transitions occur based on analysis of workload or higher-level operation.

Server systems, per [5], have characteristics that make traditional power management techniques undesirable:

- server hardware is typically provisioned for peak load, and thus exhibit high performance and power consumption;
- high availability and high bandwidth is typically provided by widespread replication of resources such as clusters of machine and disk arrays;
- server power supplies must be able to store capacity to deal with sudden spikes in load and as result exhibit high power losses.

Application servers pose an especially interesting power management problem due to intensive CPU and memory requirements combined with maintaining soft state that is not typically replicated. Intensive CPU and memory use means that traditional power management mechanisms may impose too excessive demands upon performance while non-replicated state forces state migration in order to turn application servers on and off.

One approach to dealing with the spiraling complexity in the data center is virtualization: presenting a logical grouping or subset of computing resources in a manner that provides benefits over the original configuration. Virtualization software simulates the hardware in a manner that allows a “guest” operating system to operate as if it was running on the bare hardware.

2. Motivation

Extensive research has occurred into the design and manufacture of energy and temperature efficient microprocessors. As multi-core microprocessors become more prevalent, computer architects have been exploring energy and thermal techniques required to effectively manage the energy consumption of multi-core processors [10].

Virtualization has a long history in the operating system community beginning with the IBM VM/370 commercial product [15]. In recent times, commercial systems such as VMware [28] and open source systems such as Bochs [1] have implemented virtualization systems on the Intel x86 architecture.

A common concept amongst these systems is the idea of a “hypervisor” which is the underlying piece of software responsible for managing and scheduling the different virtual machines running on a single processor. The hypervisor is responsible for managing the physical resources shared amongst the virtual machines.

The virtualization systems we have introduced so far provide “full virtualization” of the physical hardware. They must provide a complete simulation of the environment. Recently, systems such as the open source “Xen” environment [4] [2] and the commercial Parallels Desktop [23] have popularized the concept of “paravirtualization” where the guest operating system is provided with an abstraction that is similar but not identical to the underlying physical hardware. As such, guest operating systems must be aware that they are executing on a virtual machine. The advantage is that the guest

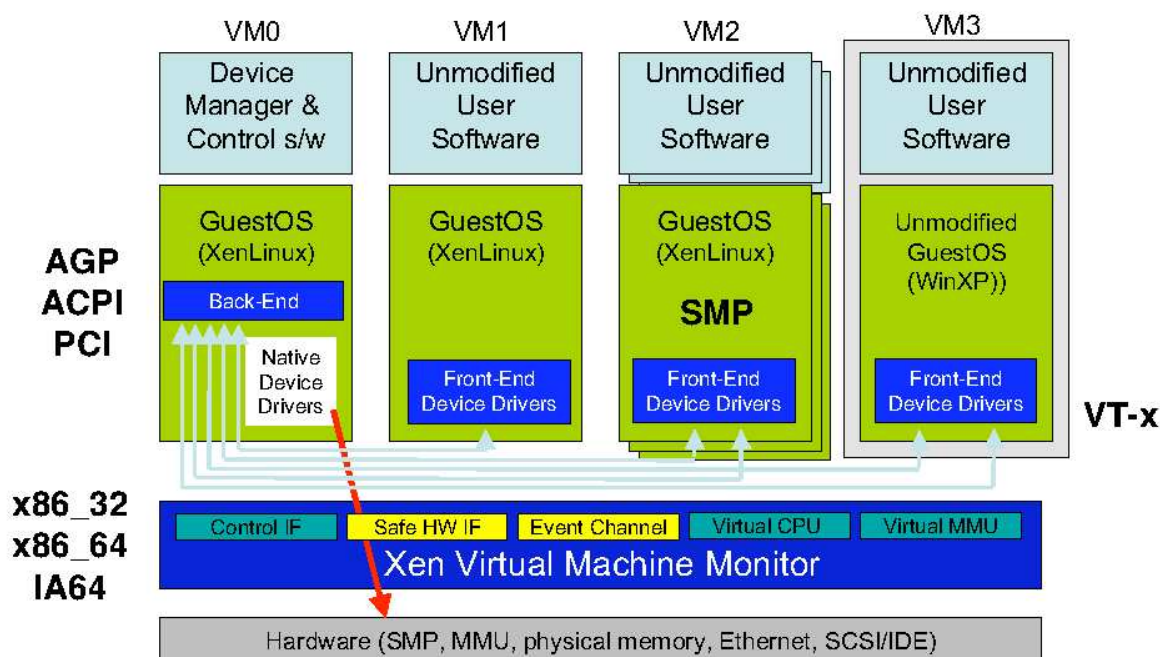


Figure 1. Xen 3.0 System Architecture [29]

can achieve performance levels approaching that of executing on the physical hardware.

In this work, we use the Xen virtualization system [4] as the starting point for our investigation into virtual machine performance. Xen was developed with the idea that it was one piece of the control systems required to build an “Internet-scale” computing infrastructure. This is important when we consider being able to provide the application developer with the view of a large collection of resources that are locally available.

The workloads typical of an application server in a computing cluster and grid often are very different than the common workload mixes in a single computer environment. Virtualization can provide high-performance computing environments with services beyond just consolidation as it allows *specialization* [21]. By consolidating physical access to the devices within the hypervisor, it allows one to tailor allocation of those resources at a much finer grain layer than what is possible with legacy operating systems while permitting the application developer to maintain some compatibility with legacy, commodity operating system.

3. Virtualization and Power Management

The conceptual model provided to the application developer by virtualization is quite seductive. From an application standpoint, all resources appear to be local and we leave it up to the hypervisor to handle the scheduling and allocation of resources. Our inspiration is the XenoServer [3] and VMPlant projects [18]. Furthermore, providing a means for the high-performance computing application to coexist with minimal interference with a commodity operating system reduces the cost and complexity of the high-performance application [21].

So, given a hypervisor capable of scheduling virtual machines amongst either multiple cores within the single die of a processor or a traditional symmetric multiprocessing system, what is the power consumption of a node in this network given the workloads expressed by a typical parallel processing benchmark in a virtualized environment.

A virtualized configuration raises some interesting questions when considering power management. First, what happens when both operating system in the hypervisor and the operating system in the virtual machine try to optimize workloads to reduce power

consumption? Second, can we achieve a better solution by just letting the hypervisor exclusively handle power management?

It should be noted that even though the focus of this research is upon minimizing the power consumption and controlling the resource allocation of computing power, similar techniques can be applied to controlling other resources within the machine (memory or disk drives, for example).

4. Design of Experiments

Statistical Design of Experiment (DoE) [22] is a methodology to efficiently determine the effects input factors have upon one or more response metrics. There are three principles in DoE: Replication, Randomization, and Blocking. Replication ensures that the results were not a fluke and reduces the error. Randomization ensures that no unintended bias is introduced by uncontrollable factors. Blocking organizes the trials so that common factors can be examined concurrently.

There are two types of input factors, qualitative and quantitative. Quantitative factors are those that can be given a value like number of wireless nodes or distance from a gateway. Qualitative factors are those which cannot be given a value as in surface terrain or routing protocol. The R-Squared value tells us how much of the variability in our data can be contributed to the input factors.

A factorial design consists several input factors are tested at different levels to and the response metrics are examined. Each input factor typically has two levels, a high and a low. Other common designs include three level factors and mixed level factors. Full factorial design means that all possible combination of factors are tested. In a 2^3 experiment, 8 trials are performed. In a half factorial design, only 4 trials would be performed. This is done when there is not enough time or resources available to complete all of the trials.

Once all of the trials have been completed, Analysis of Variance is performed to determine the factor effects and regression model. The factor effects tell us what main effects (those based on individual factors) and interaction effects (those based on the interaction between the individual factors) are important to our regression model. The regression model will give us a formula that can be used to estimate the response metric based on the input factors without running any experiments. DoE has an advantage of typical one-factor-at-a-time experiments in that it observes the effect of the interactions of the input factors. These factors are often overlooked in one—factor—at—a-time experiments and could be the most important effect in the regression model.

A set of experiments has been performed to investigate the impact of virtualization on power management. We test the hypothesis that it is more efficient to concentrate power management functionality within the hypervisor than allowing both the hypervisor and virtual machines to attempt to manage power consumption.

5. Experimental Results

First, we test whether or not varying the mix of integer and floating point instructions has an impact on the power consumption if we include the Xen hypervisor modules in the Linux kernel. We attempt to answer this question by executing a 2^2 full factorial experiment. The results of this experiment prompted further investigation with a 3^2 mixed-level factorial experiment that examined the instruction mix in more detail. The third experiment was a 2^3 experiment that examined how different combinations of instruction mix and use of virtual machines impacted the power consumed by the test hardware.

5.1 Experiment Setup and Data Collection

Our experiments take physical measurements of the power consumed by executing a collection of four benchmark programs se-

lected from the SPEC CPU2000 benchmark suite. The *gzip* and *sixtrack* benchmarks were selected due to prior work [10] finding those benchmarks drove the targeted processors to a steady state temperature.

The benchmark programs were assigned to run as independent processes either on the domain0 physical machine (i.e., the hypervisor) or on a domainU (virtual machine) controlled by the domain0 machine. The programs are executed by a driver script that runs on the domain0 machine and are started on the domainU machines using an SSH connection. We assume that the power consumption due to the overhead of this script and running the programs via SSH is minimal.

The power consumed is measured with a WattsUP [12] power meter that is connected between the AC Main and our test hardware. The power meter measures the total and average wattage, voltage, and amperage over the run of a workload. The internal memory of the power meter is cleared at the start of the run and the measures during the run are downloaded after the run completes from the meter's internal memory into a relational database [11].

Our test hardware is a Dell Poweredge 6450 configured as shown in Table 1. The operating system used is the Fedora Core 5 distribution of Linux with version 2174 of the Linux 2.6.17 kernel either with or without Xen kernel enhancements depending upon the experiment. Four identical virtual machines were created on the server.

In each experiment, we measure *power consumed* which we define as total wattage consumed over the length of run of the workload.

5.2 Is there a relationship between Xen and instruction mix?

The first of our experiments considers the relationship between power consumed, instruction mix, and whether or not the Linux kernel includes the Xen kernel modules. We address this question with the experiment shown in Table 2. This is a 2^k full factorial design with $k = 2$ factors with each factor assigned 2 values or factor levels. One combination of factors defines a design point. Each design point is replicated three times.

5.2.1 Experiment Design

The factors considered in this experiment were

IFP: The instruction Mix measures the ratio of integer instructions in the workload to the number of floating point instructions. The instruction mix was normalized to range between $-1 \leq IFP \leq 1$.

WXEN: With or without Xen indicates whether the workload was executed with an operating system kernel that contained the Xen extensions.

Note that the processes that composed the workload were all executed in the domain0 if the Xen-extended kernel was used.

5.2.2 Estimation of Effects and Allocation of Variance

The design matrix in Table 2 was analyzed using the SAS statistical software [25]. An analysis of variance (Table 3) was performed to identify the statistically significant effects and their interactions. These were combined together to build the following model:

$$TOTWATTS = 3298.167 - 112.667 * IFP$$

The effect estimates associated with this model can be found in Table 4 and a statistical analysis of the fit of the model is in Table 5. The statistics indicate we may have an issue with how well our model will predict the impact on the total wattage consumed given changes in the two inputs.

Further evidence of the lack of fit can be seen in the normal probability plot of the residuals in Figure 2. There is an assumption

Dell PowerEdge 6450	
CPU	4x700MHz Pentium III Xenon
CPU L2 cache	2MB
Memory	4GB
Internal disk	4x18GB
Network Interface Card	2x100Mbps
RAID controller	2xPERC 2/DC
Video	On-board
Height	4 rack units or 7"

Table 1. Hardware configuration




Figure 2. Normal Probability Plot for Experiment 1

in DoE [22] that the residuals (i.e., the difference between predicted value and the observed values) are normally distributed. If the residuals are plotted on the X-axis and the rank-ordered expected normal values are plotted on the Y-axis produces a graph where the values form a straight line, then one can be satisfied that the residuals are normally distributed. For this experiment, the values do not meet this criteria.

5.3 How can we confirm our conclusions about instruction mix and Xen?

The data from our previous experiment indicates that including Xen support in the kernel has statistically negligible affect upon power consumption when considering instruction mix. The instruction mix variable *IFP* is a continuous variable but we became curious about the possibility that effect on total watts consumed was not linear in nature. Additional runs were performed to expand our original experiment from a 2^2 full factorial experiment to a 3^2 full factorial experiment. The design of the expanded experiment is in Table 6.

5.3.1 Estimation of Effects and Allocation of Variance

The design matrix in Table 6 was analyzed using the SAS statistical software [25]. A second analysis of variance was performed to identify the statistically significant effects and their interactions. These were combined together to build the following model:

$$TOTWATTS = 3273.444 - 112.667 * IFP$$

A statistical analysis of the fit of the model is in Table 8. Note that both the *R - square* and *AdjustedR - square* values for the model are worse than what we saw in our first experiment. The normal probability plot of the residuals for this model is shown in Figure 3. In this case, we see a similar spread of the residuals to the previous experiment. Given this result, we see enough statistical ev-

idence to indicate that including the Xen extensions has negligible effect on power consumption when the instruction mix is varied.

5.4 How does different ways of using Xen affect power consumption?

We address this question with the experiment shown in Table 10. This is a 2^k full factorial design with $k = 3$ factors with each factor assigned 2 values or factor levels. One combination of factors defines a design point. Each design point is replicated three times.

The factors considered in this experiment were

IFP: *Instruction Mix* measures the ratio of integer instructions in the workload to the number of floating point instructions. The instruction mix was normalized to range between $-1 \leq IFP \leq 1$.

USEHYP: *Use of hypervisor* indicates whether or not one of the processes in the workload was executed in domain0.

ALLIN1: *All processes executed in one VM or spread across multiple VMs*.

5.4.1 Estimation of Effect and Allocation of Variance

The SAS statistical software [25] was used to analyze the experimental results. The main and interaction effects were analyzed using analysis of variance (Table 11) to identify those effects and interactions that are statistically significant. The resulting model includes those effects and interactions that are significant to a 95% confidence interval:

$$TOTWATTS = 3003.417 - 96.333 * IFP + 81.333 * USEHYP + 108.167 * ALLIN1$$

The statistical analysis of the fit of this predictive model to the data is in Table 12. In this case, the analysis of the fit indicates




Figure 3. Normal Probability Plot for Experiment 2




Figure 4. Normal Probability Plot of Residuals for Experiment 3

that model is a more reasonable predictor of the impact of the input factors on total wattage consumed. This conclusion is supported by the normal probability plot of the residuals in Figure 4.

Recall that the instruction mix variable *IFP* is a continuous factor varying between -1 and 1 depending upon the ratio of integer instructions to floating point instructions. As such, we can use our predictive model to build the response surface shown in Figure 5.

6. Conclusions and Future Work

We can draw the following conclusions from our experiments:

- The power consumed depending upon varying instruction mixes is not affected by having the Xen hypervisor extensions included in the system kernel.
- There is sufficient data to conclude that running processes on both the hypervisor and in the virtual machines impacts the power consumed.
- Spreading the workload across all virtual machines reduces the overall power consumed and sufficient data exists to indicate this is independent of the instruction mix.

Work has begun to investigate the impact of an energy-aware hypervisor in a single processor environment [24]. The addition of an energy-aware hypervisor capable of scheduling across multiple cores in a single processor and across multiple processors is projected to reduce the overall power consumption. Xen can be configured so that an individual virtual machine can be locked onto one CPU. Additional experiments are needed to evaluate the power profile of this configuration.

Follow-on research is needed to consider the impact of having both domain0 and domainU kernels attempting to manage the power consumption. Virtualization of these devices adds an ex-

tra level of complexity to the hypervisor that can be avoided if support for these functions can be left out of the domainU kernels.

The experimental results presented in this paper provide a baseline and testbed for evaluating the impact of other factors on power consumption of server clusters. In particular, the testbed presented in this paper will be used to evaluate the impact on power consumption of enhancements to a hypervisor's memory management scheme and virtual machine scheduling.

References

- [1] Bochs: The open source ia-32 emulation project.
- [2] Tim Abels, Puneet Dhawan, and Belasubramenaian Chandrasekaran. An overview of xen virtualization. *Dell Power Solutions*, 2005(3):109–111, August 2005.
- [3] KA Fraser and SM Hand, TL Harris, IM Leslie, and IA Prate. The xenoserver computing infrastructure. Technical report, University of Cambridge, 2003.
- [4] Paul Barham, Boris Dragovic, Keir Fraser, Steven Hand, Tim Harris, Alex Ho, Rolf Neugebauer, Ian Pratt, and Andrew Warfield. Xen and the art of virtualization. In *SOSP '03: Proceedings of the nineteenth ACM symposium on Operating systems principles*, pages 164–177, New York, NY, USA, 2003. ACM Press.
- [5] Ricardo Bianchini and Ram Rajamony. Power and energy management for server systems. *Computer*, 37(11):68–74, 2004.
- [6] H.K.F. Bjerke. Hpc virtualization with xen on itanium. Master's thesis, CERN, July 2005.
- [7] Edouard Bugnion, Scott Devine, Kinshuk Govil, and Mendel Rosenblum. Disco: running commodity operating systems on scalable multiprocessors. *ACM Trans. Comput. Syst.*, 15(4):412–447, 1997.
- [8] Stephen Childs, Brian Coghlan, David O'AoCallaghan, Geoff Quigley, and John Walsh. *Lecture Notes in Computer Science*, volume 3470, chapter Deployment of Grid Gateways Using Virtual

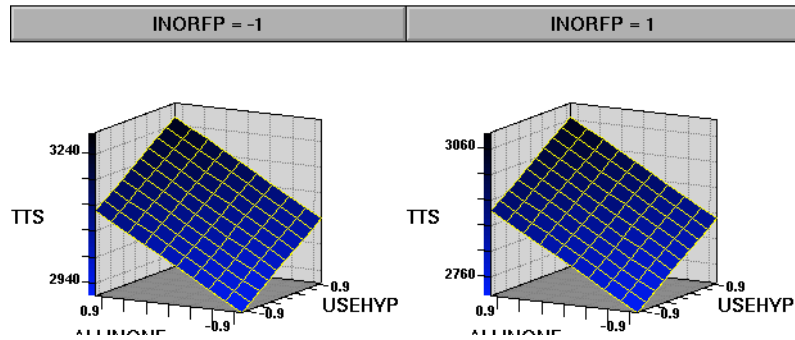


Figure 5. Response Surface for the Predictive Model in Experiment 3

- Machines, pages 761–770. Springer-Verlag, June 2005.
- [9] Stephen Childs, Brian Coghlan, David O’Callaghan, Geoff Quigley, and John Walsh. A single-computer grid gateway using virtual machines. *aina*, 01:310–315, 2005.
 - [10] James Donald and Margaret Martonosi. Techniques for multicore thermal management: Classification and new exploration. In *ISCA ’06: Proceedings of the 33rd International Symposium on Computer Architecture*, pages 78–88, Washington, DC, USA, 2006. IEEE Computer Society.
 - [11] Inc. Electronic Educational Devices. Wattsup communication protocol.
 - [12] Inc. Electronic Educational Devices, December 2006.
 - [13] Renato J. Figueiredo, Peter A. Dinda, and José A. B. Fortes. A case for grid computing on virtual machines. In *ICDCS ’03: Proceedings of the 23rd International Conference on Distributed Computing Systems*, page 550, Washington, DC, USA, 2003. IEEE Computer Society.
 - [14] Kinshuk Govil, Dan Teodosiu, Yongqiang Huang, and Mendel Rosenblum. Cellular disco: resource management using virtual clusters on shared-memory multiprocessors. *ACM Trans. Comput. Syst.*, 18(3):229–262, 2000.
 - [15] P.H. Gum. System/370 extended architecture: facilities for virtual machines. *IBM Journal of Research and Development*, 27(6):530–544, November 1983.
 - [16] Marcus Hardt and Rudiger Berlich. Xen grid site - the art of consolidation. 2005.
 - [17] Katarzyna Keahey, Ian Foster, Timothy Freeman, Xuehai Zhang, and Daniel Galron. Virtual workspaces in the grid. In *Lecture Notes in Computer Science*, volume 3648, pages 421–431. Springer-Verlag, 2005.
 - [18] Ivan Krsul, Arijit Ganguly, Jian Zhang, Jose A. B. Fortes, and Renato J. Figueiredo. Vmplants: Providing and managing virtual machine execution environments for grid computing. In *SC ’04: Proceedings of the 2004 ACM/IEEE conference on Supercomputing*, page 7, Washington, DC, USA, 2004. IEEE Computer Society.
 - [19] Charles Lefurgy, Karthick Rajamani, Freeman Rawson, Wes Felter, Michael Kistler, and Tom W. Keller. Energy management for commercial servers. *Computer*, 36(12):39–48, 2003.
 - [20] Bin Lin and Peter A. Dinda. Vsched: Mixing batch and interactive virtual machines using periodic real-time scheduling. In *SC ’05: Proceedings of the 2005 ACM/IEEE conference on Supercomputing*, page 8, Washington, DC, USA, 2005. IEEE Computer Society.
 - [21] Mark Mergen, Volkmar Uhlig, Orran Krieger, and Jimi Xenidis. Virtualization for high-performance computing. *ACM SIGOPS Operating Systems Review*, 40(2), April 2006.
 - [22] Douglas C. Montgomery. *Design and Analysis of Experiments*. John Wiley & Sons, Inc, 11 River Street, Hoboken, NJ 07030, 6th edition, 2005.
 - [23] Inc Parallels. Parallels workstation.
 - [24] Marcus Reinhardt. Cooperative, energy-aware scheduling of virtual machines. Master’s thesis, University of Karlsruhe, August 2005.
 - [25] Inc. SAS, December 2006.
 - [26] Srikanth Sundararajan, Hariprasad Nellitheertha, Subhabrata Bhattacharya, and Neel Arurkar. Nova: An approach to on-demand virtual execution environments for grids. *ccgrid*, 0:544–547, 2006.
 - [27] Rich Uhlig, Gil Neiger, Dion Rodgers, Amy L. Santoni, Fernando C. M. Martins, Andrew V. Anderson, Steven M. Bennett, Alain Kagi, Felix H. Leung, and Larry Smith. Intel virtualization technology. *Computer*, 38(5):48–56, 2005.
 - [28] VMware. Vmware esx server 2 architecture and performance implications. Technical report, VMware Corporation, 2005.
 - [29] Inc. XenSource. Xensource,inc.
 - [30] X. Zhang, T. Freeman, K. Keahey, and I. Foster adn D. Scheftner. Virtual clusters for grid communities. 2006.

<i>RUN</i>	<i>IFP</i>	<i>WXEN</i>	<i>TOTWATTS</i>
2	1	-1	3066
6	1	-1	3193
10	1	-1	3198
4	1	1	3225
8	1	1	3235
12	1	1	3196
1	-1	-1	3450
5	-1	-1	3454
9	-1	-1	3450
3	-1	1	3458
7	-1	1	3196
11	-1	1	3457

Table 2. 2^2 Design Matrix and Experimental Output for Experiment 1

Source	<i>DF</i>	Master Model				<i>DF</i>	Predictive Model			
		<i>MS</i>	<i>MS</i>	<i>F</i>	<i>Pr > F</i>		<i>SS</i>	<i>MS</i>	<i>F</i>	<i>Pr > F</i>
<i>IFP</i>	1	152325.3	152325.3	21.15164	0.001758	1	152325.3	152325.3	20.5694	0.001082
<i>WXEN</i>	1	161.3333	161.3333	0.022402	0.884726					
<i>IFP * WXEN</i>	1	16280.33	16280.33	2.26066	0.171106					
<i>Model</i>	3	168767	56255.67	7.811569	0.00921	1	152325.3	152325.3	20.5694	0.001082
<i>Error</i>	57612.67	7201.583			10	74054.33	7405.433			
<i>Total</i>	11	226379.7				11	226379.7			

Table 3. ANOVA Analysis for Experiment 1

Term	Master Model				Predictive Model			
	Estimate	Std Err	t	<i>Pr > t </i>	Estimate	Std Err	t	<i>Pr > t </i>
<i>IFP</i>	-225.333	48.9958	-4.599	0.00176	-225.330	49.685	-4.535	0.001082
<i>WXEN</i>	-7.333	48.995	-0.1500	0.88473				
<i>IFP * WXEN</i>	73.6667	48.99518	1.504	0.17110				

Table 4. Effect Estimates for Experiment 1

	Master Model	Predictive Model
<i>RMSE</i>	84.8621	86.0548
<i>R – square</i>	74.55%	67.25%
<i>Adjusted R – square</i>	65.01%	64.02%
<i>Coefficient of variation</i>	2.573	2.609

Table 5. Fit Statistics for Experiment 1

<i>RUN</i>	<i>IFP</i>	<i>WXEN</i>	<i>TOTWATTS</i>
3	1	-1	3066
9	1	-1	3193
15	1	-1	3198
6	1	1	3225
12	1	1	3235
18	1	1	3196
28	0	-1	3494
8	0	-1	3187
14	0	-1	3179
5	0	1	3142
11	0	1	3158
17	0	1	3184
1	-1	-1	3450
7	-1	-1	3454
13	-1	-1	3450
4	-1	1	3458
10	-1	1	3196
16	-1	1	3457

Table 6. 3^2 Design Matrix and Experimental Output for Experiment 2

Source	DF	MS	Master Model			DF	SS	Predictive Model		
			MS	F	Pr > F			MS	F	Pr > F
IFP	1	152325.300	152325.300	13.0392	0.00257	1	152325.300	152325.300	13.172	0.002255
WXEN	1	9800.000	9800	0.839	0.374					
Model	2	162125.300	81062.670	6.939065	0.00735	1	152325.300	152325.300	13.1719	0.00226
Error	15	175231.100	11682.070			16	185031.100	11567.440		
(Lack of Fit)	3	52207.110	17402.370	1.975	0.220	1	22002.780	22002.780	2.0244	0.1753
Total	17	337356.400				17	337356.400			

Table 7. ANOVA Analysis for Experiment 2

	Master Model	Predictive Model
RMSE	108.084	107.538
R – square	48.06%	45.15%
Adjusted R – square	41.13%	41.72%
Coefficient of variation	3.301832	3.285167

Table 8. Fit Statistics for Experiment 2

Term	Estimate	Master Model			Estimate	Predictive Model		
		Std Err	t	Pr > t		Std Err	t	Pr > t
IFP	-192.667	50.7267	-3.798	0.00144	-192.667	50.691	-3.8008	0.00112
WXEN	-23.333	25.475	-0.916	0.374				

Table 9. Effect Estimate for Experiment 2

RUN	IFP	USEHYP	ALLIN1	TOTWATTS
8	1	1	1	3071
16	1	1	1	3051
24	1	1	1	3032
4	1	1	-1	2862
12	1	1	-1	2890
20	1	1	-1	3082
6	1	-1	1	2929
14	1	-1	1	2835
22	1	-1	1	2965
2	1	-1	-1	2732
10	1	-1	-1	2704
18	1	-1	-1	2732
7	-1	1	1	3287
15	-1	1	1	3286
23	-1	1	1	3282
3	-1	1	-1	3063
11	-1	1	-1	2847
19	-1	1	-1	3264
5	-1	-1	1	3464
13	-1	-1	1	3069
21	-1	-1	1	3068
1	-1	-1	-1	2781
9	-1	-1	-1	2785
17	-1	-1	-1	3001

Table 10. 2^3 Design Matrix and Experimental Output for Experiment 3

Source	DF	MS	Master Model			DF	SS	Predictive Model		
			MS	F	Pr > F			MS	F	Pr > F
IFP	1	222722.700	222722.700	14.426	0.00144	1	222722.700	222722.700	14.426	0.00144
USEHYP	1	158762.700	158763.700	10.283	0.00517	1	158762.700	158762.700	10.297	0.004405
ALLIN1	1	280800.7	280800.7	18.188	0.00523	1	280800.7	280800.700	18.2129	0.000376
IFP * USEHYP	1	2204.167	2204.167	0.143	0.710					
IRFP * ALLIN1	1	28981.500	28981.500	1.877	0.188					
USEHYP * ALLIN1	1	14701.500	1471.500	0.952	0.343					
Model	6	708173.200	118028.900	7.645	0.00427	3	662286	220762	14.310	0.00001
Error	17	262466.700	15439.220			20	308353.8	154417.690		
(Lack of fit)	1	522.667	522.667	0.0319	0.860	4	46409.83	11602.460	0.709	0.598
(Pure Error)	16	2619441.000	16371.0005			16	261944	16371.5		
Total	23	970639.8				23	970639.8			

Table 11. ANOVA Analysis for Experiment 3

	Master Model	Predictive Model
RMSE	124.2546	124.168
R-square	72.96%	68.23%
Adjusted R-square	63.42%	63.47%
Coefficient of variation	4.137	4.134

Table 12. Fit Statistics for Experiment 3

Term	Master Model				Predictive Model			
	Estimate	Std Err	t	Pr > t	Estimate	Std Err	t	Pr > t
IFP	-192.667	50.728	-3.798	0.00144	-192.667	50.691	-3.801	0.00112
USEHYP	16.667	50.727	3.207	0.00517	162.667	50.691	3.209	0.00441
ALLIN1	216.333	50.727	4.267	0.000523	216.333	50.691	4.268	0.000376
IFP * USEHYP	19.1667	50.727	0.378	0.710				
INORFP * ALLIN1	-69.5	50.7274	-1.370	0.188				
USEHYP * ALLIN1	-49.5	50.7274	-0.9762	0.343				

Table 13. Effect Estimate for Experiment 3