Thermally-aware Load Balancing Using Predictive Energy Consumption Models in High-Performance Systems

Adam Lewis, Soumik Ghosh, and N.-F. Tzeng

Center for Advanced Computer Studies,University of Louisiana, Lafayette, Louisiana 70504
{awlewis,sxg5317,tzeng}@cacs.louisiana.edu

Abstract

Power management techniques developed for mobile and desktop computers are not always appropriate for high-performance computing systems because resources on these systems tend be fully utilized.As the server workload increases, so does the thermal stress on the server processors. Modern processors crudely manage this problem through Dynamic Thermal Management (DTM) where the processor monitors the die temperature and applies Dynamic Voltage/Frequency Scaling (DVFS) to slow the processor and reducing energy consumption and thermal impact. However, DVFS methods entail a heavy cost from both a performance and reliability standpoint.

This paper introduces policy-based, thermally-aware load balancing scheduler that utilizes a full-system thermal and energy model to proactively predict and minimize hazardous thermal effects. The dataflow-based systemwide model estimates the energy consumption and thermal impact of an application during its execution. The knowledge of the run-time predictive thermal status of the processor allows the scheduler time to migrate threads and applications to cooler, underutilized cores with minimal performance impact. The model uses key micro-architectural event counters and operating system kernel statistics to keep track of workload related activity within the processor and associated peripherals.

The scheduler uses the die temperature estimates to adjust the load balancing policies to minimize the number of DTM events. The analytical model and the scheduler performance is demonstrated through empirical evaluation to outperform previous approaches. This is shown by extending an existing power-aware scheduler that is part of the OpenSolaris operating system executed on an AMD Opteron processor. A subset of the SPEC CPU2006 benchmark suite is used to simulate application server workloads and evaluate scheduler performance.

1 Introduction

Power management techniques developed for mobile and desktop computers are not always appropriate for real-time and high-performance computing applications as these techniques take advantage of slack periods that occur within the application workloads. High-performance computing workloads result in the application server being fully utilized with few resources available for reactive scheduling.

As the workload increases on the server, so does the thermal stress placed on the processor with the resulting probability of damage to the machine. Modern processors crudely manage this problem through Dynamic Thermal Management (DTM) where the processor monitors the die temperature and dynamically adjusts the processor voltage and frequency (DVFS) to slow down the processor. However, DVFS has a significant negative impact on application performance. As such, it such it becomes necessary to find power management techniques that are pro-active in nature rather than reactive.

In this work, we introduce a technique for thermal management that minimizes server energy consumption by adjusting the allocation of thread groups to available cores so as to produce the least contention for memory accesses amongst thread groups. This must be balanced against the performance benefits that are gained through cache affinity; by allocating thread groups to last utilzied cores as as to gain performance by taking advantage of data still existing in cache will lead to thermal issues through more active cache access.

In summary, the contributions of this work are:

- We show how a statistical full-system model of energy consumption of server blades can be extended to model the thermal behavior of those blades.

- We use this model as a predictor of processor thermal behavior in a policy-based thermal scheduler.

- We evaluate our model and scheduler by extending an existing power-aware scheduler running in OpenSolaris on an AMD Opteron processor.

2 Related Work

Previous attempts to implement effective thermal management in processors have taken two distinct approaches: hardware based implementations of Dynamic Thermal Management (DTM) combined with Dynamic Voltage-Frequency Sacaling (DVFS) or operating system based multiprocessor or multi-core thread migration approaches.

Commercial processors crudely regulate thermal stress using DTM by slowing down the processor using DVFS whenever the processor temperature exceeds a certain threshold. A categorization of the different thermal management techniques based on this approach can be found in [10]. In this work, the authors claimed the best approach to using DTM/DVFS includes a combination of a control-theoretic distributed DVFS system and a sensor-based migration policy. The issue with

this approach is the granularity of DVFS; existing commercial systems do not allow for independent frequency scaling of cores and threads. A related approach is the capability of turning off the power to unused cores in the processor. However, use of this capability for power management can lead to reduced reliability in the processor [26][9].

Furthermore, there is a significant cost to the use of DVFS as a method to avoid thermal emergencies. Bircher *et.al.* [5] examined the impact of dynamic power adaptions on the performance and power envelope of a Quad-Core AMD Opteron processor. In this work, it was found that significant performance loss is entailed due to slow transitions between different levels of DVFS. In addition, indirect effects due to shared, power-managed resources such as cache and memory can greatly reduce performance; this is a pronounced problem for memory-bound workloads typical of server workloads.

The idea of migration of work for energy savings and thermal management has a long history in the SMP, SMT, and CMP environments [30]. Static methods been applied in attempts to solve this problem. In Coskun `et. al.`[8], integer linear programming was used to obtain a task schedule that met real-time deadlines while attempting to minimize hotspots and spatial temperature differentials across the die.

Examples of dynamic methods to solve this problem incude Heat-and-Run [13],HybDTM [18], and ThreshHot [29]. Heat-and-Run proposed to distribute work amongst available cores until the DTM events occur and then migrate threads from the overheated cores to other non-heated coraes. HybDTM combined DTM techniques with a thread migration strategy that reduces the thread priority of jobs on cores that are running hot. ThreshHot uses an on-line temperature estimator to determine in what order threads should be scheduled onto cores. In all three cases, these techniques utilize data from hardware performance counters and hardware temperature sensors to guide the scheduling decision.

A related approach to migration is to control the CPU time slice allocation. In Bellosa *et.al.*[3], it was proposed to modify the process scheduler to allocate time slices as indicated by the contribution of each task to the system power consumption and the current temperature of the processor. A variation on this scheme was proposed in [20] where system level compiler support was used to insert run-time profiling code into applications to provide hints to the thermal intensity of a task. In Merkel `et.al.` [22], a scheduling policy was proposed that sorts the tasks in each core's run queue by memory intensity so as to schedule memory-bound tasks at slower frequencies.

2.1 Power and Energy Models

The foundation of these techniques is a power and energy model of the underlying processor. These models can be classified into two broad categories: detailed analytical power models and high-level blackbox models [25]. Analytical models use detailed knowledge of the underlying hardware to either simulate the energy consumption or directly measure energy comsumption at the hardware level. Simulation can provide detailed analysis and breakdown of energy consumption; however, they are slow and do not scale well to realistic applications and large data sets. Simulation-based models do not fit well into scensarios where dynamic optimization of application performance is required [11].

High level black-box models sacrifice some accuracy by avoiding extensive detailed knowledge of the underlying hardware. At the processor level, [6], and [3] created power models that linearly correlated power consumption with performance counters. Models have been built for the processor, storage devices, single systems, and groups of systems in data centers. These models have the advantage of being simple, fast and low-overhead but do not model the full-system power consumption.

Measurement-based models attempt to collate power measurements with hardware and software performance metrics. Two distinct classes of metrics have been used in these models: processor performance counters and operating system performance metrics. Processor performance counters are hardware registers that can be configured to count various microarchitectural events such as branch mispredictions and cache misses. In general, the number of countable events exceed the number of available registers. As result, models that use these counters time-multiplex different sets of events on the available registers. While this allows for more events to be monitored, it results in increased overhead and lower accuracy due to sampling issues [11][24].

Attempts have been made to reconcile these approaches by attempting to map programs phases to events [16]. The most common technique used to associate PeCs and/or operating systems metrics to energy consumption use linear regression models to the collected metrics to the energy consumed during the execution of a program [6][11][17][4].

In server envrionments, it has been shown that full-system models using operating system CPU utiliziation can be highly accurate [12]. Others have used similar approaches [14] to develop linear models for energy-aware server consolidation in clusters. Full-system models such as MANTIS [11][24] relate usage information to the power of the entire system rather than an individual component. The accuracy and portability of full system power models is considered in [25]. This analysis indicated that to ensure reasonable accuracy across machines and workload required a model based on a combination of both PeCs and operating system metrics.

2.2   Thread Migration and Load Balancing

Recent work has concentrated on building proactive thermal control mechansims by combining thread migration with a predictive model of workload impact on thermal changes and energy consumption. Rangan, Wei, and Brooks [23] introduce the idea of thread motion where, given similar processing cores operating at different voltage/frequencey level, workload is migrated to cores operating at a higher or lower voltage/frequency level depending upon workload demand.

The work most similar in spirit to ours is that of Coskun `et.al.` [9] where an autoregressive moving average model of time series temperature sensor data is used to modify the load balancing decisions made by the operating system dispatcher. The predictive model used in this work attempts to adapt their model to workload based on evaluation of prediction error. Our work avoids this expensive recalibration by incoporating evaluation of impact of architectural events into the model.

3   Predictive Thermal Load Balancing

```
Where are hotspots located on the processor?
Switching, buses, and cache activity
Performance load balancing encourages hotspots
Predict changes in temperature based on this activity.
Provide load balancing policies that move work away from areas
```

Our scheduler is an enhancement of the existing thread scheduling infrastructure found in modern operating systems.

4   Energy Models and Performance Counters

The model used predicting energy consumption in our scheduler is based upon the model introduced in [19]. This model provides a system-wide view of the energy consumption by using hardware performance counters to relate system power consmption to its overall thermal envelope.

Modern processor cores support between 30 to 500 PeCs (depending upon the processor) but only permit 2 to 18 of these counters to be read at one time (again depending upon processor). Furthermore, certain combinations of PeCs may not be permitted to be collected at the same time. Using PeCs to gain insight into architectural events (in our case bus transactions)is complicated by the fact that the types of events, number of events, and use cases varies widely, not only across architectures, but across systems sharing the same Instruction Set Architecture (ISA). For instance, the Intel and AMD implementations of performance counters have very little in common in spite of the processor families using the same ISA [27][1].

As an example, to derive the thermal impact of application workloads on cache behavior, the scheduler needs to collect nine PeCs to calculate shared and unshared cache activity on the AMD Opteron processor: $RetiredInstructions$, $DCAccesses$, $DCRefillsL2$, $DCRefillsFromSystem$, $ICFetches$, $ICrefillsFromL2$, $ICRefillsFromSystem$, $L2RequestsTLB$, and $L2MissesTLB$. The model must calculate the L2 miss rate and L2 miss ratio using the following set of formulas with the PeCs as inputs:

$$L2MissRate = L2misses/RetiredInstructions$$
$$L2MissRatio = L2misses/L2Requests$$

These measures indicate cache activity which our model uses as an estimator for energy consumption and potential for a DTM event.

Note that we have to collect nine different PeCs even though the processor can only collect four of these counters at a time. For real-time situations, time multiplexing is the most popular solution to this problem [2]. In this approach the performance

counters are reconfigured for different sets of counter events at regular time intervals. However, time multiplexing introduces issues with reconfiguration overhead and time alignment of samples.

4.1 Thermal Extensions to the Model

There are two major metrics of interest for the thermal workload of a multicore processor. The first is the total time of execution (denoted by L for length) of a certain application (denoted by A),

$$L(A(p_i), D_A, t)$$

where application A has p processes, each associated with a data set of size d_i in a single chip. D_A is the total data associated with application A, and $D_A = \sum_{i=1}^{p} d_i$. We assume that the activities are taking place in a staging area which contains the main and virtual memory operating spaces, as well as the processor with its cores and their associated caches and shared cache.

This time of execution measurement includes both computation time, and the time to move the data for the problem from the staging area (peripherals off the chip like DRAM and HDD) to a computation or operation area (on the chip such as the caches and the cores).

The second major metric of interest is the energy consumption or energy workload of an application, $U(A(p_i), d_i)$. For each application A and problem size D_A, a measure of the workload, $W(A(p_i), d_i)$, is defined in a data-operation dependent and system-independent way. W contains two components, one being the operations count that is performed by the computational core, and the other is the communication operations in transferring data and instructions and data coherency and book-keeping operations. These are measured in terms of the number of bytes operated ON, or number of bytes transferred. Thus the energy workload of an application A operating on a data set D_A can be expressed as :

$$U(A(p_i), d_i) = \lim_{n \to k_e} n \times W(A(p_i), d_i) \times L_n(A(p_i), d_i, t)$$

where $L_n(k, d_i)$ is the total time to execute n applications using the chip. The term k_e is the total number of applications that can be excuted with the associated length of time for L_n, at which point a "thermal event" will occur causing the applications and the system to catastrophically fail, or shut down.

It is easy to see that the above term is energy consumption of the system till a thermal event occurs. In order to relate the energy expenditure of the system while running applications, to teh corresponding joule heating, we define the term "Thermal equivalent of Application" (TEA), which is defined as the electrical work converted to heat in running an application and is measured in terms of die tmperature change and ambient temperature change of the system. Thus for the application A we express TEA as :

$$\Theta_A(A(p_i), d_i, T, t) = \frac{U(A(p_i), d_i)}{\lim_{T \to T_{th}} J_e \times (T - T_{nominal})}$$

$$= \frac{\lim_{n \to k_e} n \times W(A(p_i), d_i) \times L_n(A(p_i), d_i, t)}{\lim_{T \to T_{th}} J_e \times (T - T_{nominal})}$$

In these expressions, T_{th} refers to the threshold temperature at which a DTM triggered event will occur. $T_{nominal}$ refers to the nominal temperature as reported by the DTM counters / registers , when only the operating system in operating and no application is being run. The term J_e is the "electrical equivalent of heat" for the chip, which reflects the *informational entropy* of the system associated with processing the bits application $A(p_i)$ computes and communicates, as well as the black body thermal properties of the chip packaging and the cooling mechanisms around the chip. TEA thus is a dimensionless quantity , both denominator and numerator being expressing of work done or energy consumed in finishing a task.

For managing the thermal envelope of applications on server systems as well as embedded systems, we are interested in the thermal efficiency of the operation, that is, the thermal cost of taking an application to completion. In general the efficiency $\eta(A(p_i), d_i, T)$ is defined as

$$\eta(A(p_i), d_i, T, t) = \frac{\Theta_A(A(p_i), d_i, T, t)}{\Theta_A(A_e(p_i), d_i, T_{me}, t_e)}$$

where $T_{(me)}$ is the maximum temperature till which the core will carry over till a DTM triggered event occurs. A_e refers to the application whose energy consumption has caused the DTM triggered event to take place. T_e is the execution time of application A_e. Thus $\eta(A(p_i), d_i, T)$ is a measure of the "thermal efficiency of the application", which implies how much an application affects temperature change without compromising it's throughput and/or leads to a thermal event. Thus the definition of η is linked to the definition of the thermal and energy workload.

An important metric finally is the achieved performance per unit power consumed by the chip,

$$C_\theta(A(p_i), d_i, T, t) = \frac{\Theta_A(A(p_i), d_i, T, t)}{P(A(p_i), d_i)}$$

$$= \frac{\Theta_A(A(p_i), d_i, T, t)}{\sum^{chip,DRAM,HT,HDD} \int_{t=0}^{t=L_A} v(t)i(t)dt}$$

where $P(A(p_i), d_i)$ is the overall power consumed during the application lifetime. the quantity $\int_{t=0}^{t=L_A} v(t)i(t)dt$ is the total power consumed by a single physical component (processor, DRAM units, HDD, HT or FSB) during the length of the application L_A, with $v(t)$ and $i(t)$ being the instantaneous voltage and currents respectively.

This normalized quantity C_θ gives some indication of the "cost" of executing the benchmark on the given chip.

The final optimization function could be thought of as :

$$\frac{\partial^2 C_\theta(A(p_i), d_i, T, t)}{\partial T \partial t} = \frac{\partial^2}{\partial T \partial t} \frac{\lim\limits_{n \to k_e} n \times W(A(p_i), d_i) \times L_n(A(p_i), d_i, t)}{\lim\limits_{T \to T_{th}} J_e \times (T - T_{nominal})} \times \frac{1}{\sum\limits_{chip, DRAM, HT, HDD} \int_{t=0}^{t=L_A} v(t) i(t) dt}$$

```
Need to work mapping of model to physical quantities measured (PeCs,
other metrics)
```

```
This will be a table plus additional formulas
```
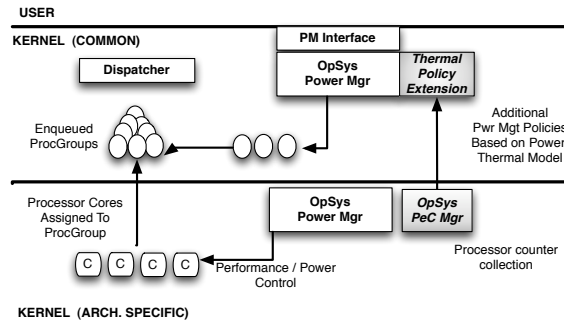

Figure 1. PAD thermal enhancement.

5 Design and Implementation

Our scheduler extends the existing dispatcher and power management infrastructure in the operating system. It is the role of the kernel dispatcher to manage the placement of threads in a dispatch queue, decide which thread to run on a processor, and manage the movement of threads to and from processors. At the fundamental level, the kernel dispatcher is a queue management system that must perform four core functions [21]: (1) queue management, (2) thread selection, (3) processor selection, and (4) context switching.

5.1 Solaris Power-Aware Dispatcher

OpenSolaris defines locality groups (lgroups) to provide a locality-oriented mechanism to represent a collection of CPUs and memory resources that are within some latency of each other. Lgroups are hierarchical and are created automatically by Solaris based on the systems conguration and different levels of locality. The system places a new thread into a home lgroup that is based on load averages although applications can give a thread a different home lgroup. Lgroups help to control where threads and memory are allocated. When a thread is scheduled to run, the thread is assigned to the available CPU nearest to the home lgroup, and memory is gotten either from the home lgroup or some parent lgroup.
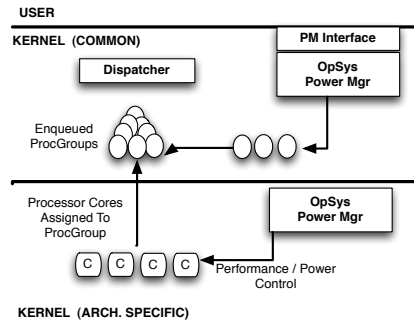
Figure 2. Solaris Power-Aware Dispatcher design.

Inside the Solaris kernel, logical CPUs will belong to one or more CPU partitions. An overall system CPU partition is created at system initialization. At the user level, a logical CPU will be a member of one or more processor groups. Processor groups are implemented inside the kernel as CPU partitions. The Solaris Power-Aware Dispatcher (PAD) (Figure 2) [28] extended the Solaris kernel dispatcher to enumerate logical CPUs that represent active and idle power domains. The concept of processor groups has been extended to include the concept of logical CPUs sharing one of two power domains: active and idle. At system start-up, the dispatcher will query the operating system's power manager to determine if a logical CPU is a member of either of these domains. If so, the CPU is added to the appropriate processor group. As the workload progresses, the power manager tracks utilization in each power domain. The power manager utilizes this information to decide how to change power states for each CPU. This capability allows the dispatcher to attempt to concentrate light workloads on a smaller number of logical CPUs and make more resources available for scheduling.
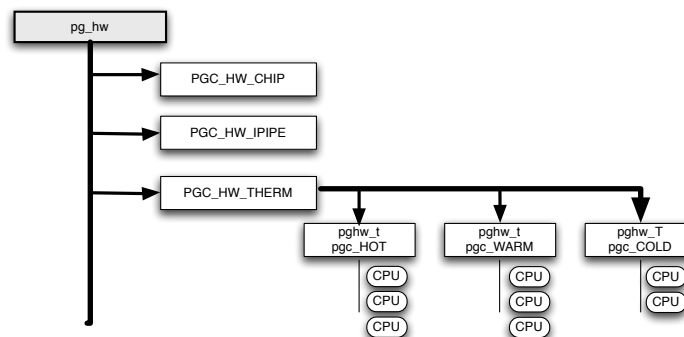

Figure 3. Thermal processor groups.

5.2 Extending Processor Groups for Thermal Management

The concept of processor group is extended to include three new processor domains organized by the temperature of the logical CPUs in this time period: hot (90% of DTM temperature), warm (between 75% and 90% of DTM temperature), and

cold (less than 75% of the DTM temperature). As the temperature of a logical CPU increases, the processor core may be assigned to a different domain based on how close the temperature is to causing a DTM event. The movement of logical CPUs


Figure 4. Thermal predictor inputs.

between processor groups as temperature increases is managed by a Thermal Predictor. The Thermal Predictor collects the values of the PeCs, temperature readings from the logical CPUS, and utilization data from the dispatcher to decide whether a processor will move towards the DTM temperature. In this way we can anticipate the DTM event and prevent its occurance. The dispatcher is responsible for deciding where to insert threads into run queues; i.e., when and where a thread will next
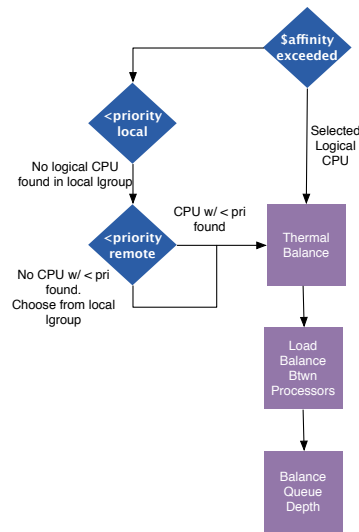

Figure 5. Thread Queue Insertion

execute. The steps in this process are illustrated in Figure 5. First, the dispatcher checks to see if the logical CPU where the thread was running satisifies the cache affinity critieria (the default for OpenSolaris is the thread was last running within the past 3 cycles). If that test fails, it tries to find the logical CPU in the local lgroup running the lowest priority thread. If no thread is found in the local lgroup, then the dispatcher will check through known remote lgroups. If no logical CPU is found in either case, it will choose a logical CPU in the local lgroup.

At this point, the system will attempt to load balance the workload.

Two possibilities exist in the current system: equal balance across all physical processors in the system and coalese work

Table 1. SPEC CPU2006 Benchmarks Used in Callibration

Benchmark		Type	Use
perlbench	C	Integer	PERL Programming Language
bzip2	C	Integer	Compression
mcf	C	Integer	Combinatorial Optimization
omnetpp	C++	Integer	Discrete Event Simulation
gromacs	C/Fortran	Floating Point	Biochemistry/Molecular Dynamics
cacstusADM	C/Fortran	Floating Point	Physics/General RElativity
leslie3d	Fortran	Floating Point	Fluid Dynamics
lbm	C	Floating Point	Fluid Dynamics

onto one chip or the other. The first case is traditional example of load balancing while the second case is attempting to free up as many resources as possible to provide opportunities for power management to find logical units to shut down. We add a third sort of load balancing: temperature balance. In this case, the dispatcher will attempt to load logical CPUs in the "COLD" processor group first, followed by the "WARM" processor group, and finally the "HOT" group.

6 Experimental Evaluation

6.1 Workloads

Typical server workloads are represented in our environment by the use of a selected series of benchmarks taken from the SPEC CPU2006 benchmark suite [15] and the SPEC SPECpower benchmark (Table 1). A multiprogrammed server environment was simulated in our test environment by exhaustively running pairs of benchmarks from the set of SPEC CPU2006 benchmarks.


Figure 6. Thermal-aware scheduling test fixture.

6.2 Measurement Environment

The hardware used to evaluate our model is described in Table 6.2. The power consumed is measured with a Yokogawa WT210 power analyzer [7] connected between the AC Main and the SUT. The power meter measures the total and average wattage, voltage, and amperage over the run of a workload. The internal memory of the power meter is cleared at the start of

Table 2. Test Hardware Configuration

	Sun Fire 2200
CPU	2 AMD Opteron
CPU L2 cache	2x2MB
Memory	8GB
Internal disk	2060GB
Network Interface Card	2x1000Mbps
Video	On-board
Height	1 rack unit

the run and the measures collected during the run are downloaded after the run completes from the meter's internal memory into a spreadsheet.

6.3 Benchmark Baseline Classification

6.4 Thermal Scheduling Results

7 Discussions

8 Conclusions and Future Work

References

[1] AMD. *BIOS and Kernel Developer's Guide for AMD Athlon 64 and AMD Opteron Processors*, 26094 rev 3.30 edition, February 2006.

[2] R. Azimi, M. Stumm, and R. W. Wisniewski. Online Performance Analysis by Statistical Sampling of Microprocessor Performance Counters. In *ICS '05: Proceedings of the 19th Annual Intl. Conf. on Supercomputing*, pages 101–110, New York, NY, USA, 2005. ACM.

[3] F. Bellosa, A. Weissel, M. Waitz, and S. Kellner. Event-driven energy accounting for dynamic thermal management. *Proceedings of the Workshop on Compilers and Operating Systems for Low Power (COLP'03)*, 2003.

[4] W. Bircher and L. John. Complete System Power Estimation: A Trickle-Down Approach Based on Performance Events. *Ispass*, 0:158–168, 2007.

[5] W. L. Bircher and L. K. John. Analysis of dynamic power management on multi-core processors. In *ICS '08: Proceedings of the 22nd annual international conference on Supercomputing*, pages 327–338, New York, NY, USA, 2008. ACM.

[6] G. Contreras and M. Martonosi. Power Prediction for Intel XScale®Processors Using Performance Monitoring Unit Events. In *ISLPED '05: Proc. of the 2005 Intl. Symp. on Low Power Electronics and Design*, pages 221–226, New York, NY, USA, 2005. ACM.

[7] Y. E. Corporation. *WT210/WT230 Digital Power Meter User's Manual*. Yokogawa Electric Corporation, 4th ed. edition.

[8] A. K. Coskun, T. S. Rosing, K. A. Whisnant, and K. C. Gross. Temperature-aware mpsoc scheduling for reducing hot spots and gradients. In *ASP-DAC '08: Proceedings of the 2008 conference on Asia and South Pacific design automation*, pages 49–54, Los Alamitos, CA, USA, 2008. IEEE Computer Society Press.

[9] A. K. Coskun, T. v. Rosing, K. A. Whisnant, and K. C. Gross. Static and dynamic temperature-aware scheduling for multiprocessor socs. *IEEE Trans. Very Large Scale Integr. Syst.*, 16(9):1127–1140, 2008.

[10] J. Donald and M. Martonosi. Techniques for Multicore Thermal Management: Classification and New Exploration. In *ISCA '06: Proc. of the 33rd Intl. Symp. on Computer Architecture*, pages 78–88, Washington, DC, USA, 2006. IEEE Computer Society.

[11] D. Economou, S. Rivoire, C. Kozyrakis, and P. Ranganathan. Full-System Power Analysis and Modeling for Server Environments. In *Workshop on Modeling Benchmarking and Simulation (MOBS) at ISCA*, 2006.

[12] X. Fan, W.-D. Weber, and L. A. Barroso. Power provisioning for a warehouse-sized computer. In *ISCA '07: Proc. of the 34th Intl. Symp. on Computer architecture*, pages 13–23, New York, NY, USA, 2007. ACM.

[13] M. Gomaa, M. D. Powell, and T. N. Vijaykumar. Heat-and-run: leveraging smt and cmp to manage power density through the operating system. *SIGOPS Oper. Syst. Rev.*, 38(5):260–270, 2004.

[14] T. Heath, B. Diniz, E. V. Carrera, W. M. Jr., and R. Bianchini. Energy Conservation in Heterogeneous Server Clusters. In *PPoPP '05: Proc. of the 10th ACM SIGPLAN Symp. on Principles and Practice of Parallel Programming*, pages 186–195, New York, NY, USA, 2005. ACM.

[15] J. L. Henning. SPEC CPU2006 Benchmark Descriptions. *SIGARCH Comp. Archit. News*, 34(4):1–17, 2006.

[16] C. Isci and M. Martonosi. Phase Characterization for Power: Evaluating Control-flow-based and Event-counter-based Techniques. *The 12th Intl. Symp. on High-Performance Computer Architecture*, pages 121–132, 11-15 Feb. 2006.

[17] C. Isci and M. Martonosi. Runtime Power Monitoring in High-end Processors: Methodology and Empirical Data. *MICRO-36: Proc. 36th IEEE/ACM Intl. Symp. on Microarchitecture*, pages 93–104, 3-5 Dec. 2003.

[18] A. Kumar, L. Shang, L.-S. Peh, and N. K. Jha. Hybdtm: a coordinated hardware-software approach for dynamic thermal management. In *DAC '06: Proceedings of the 43rd annual conference on Design automation*, pages 548–553, New York, NY, USA, 2006. ACM.

[19] A. Lewis, S. Ghosh, and N.-F. Tzeng. Run-time energy consumption estimation based on workload in server systems. In *Proc. of the 2008 Workshop on Power Aware Computing and Systems (Hotpower'08)*, 2008.

[20] D. Li, H.-C. Chang, H. Pyla, and K. Cameron. System-level, thermal-aware, fully-loaded process scheduling. pages 1–7, April 2008.

[21] R. McDougall and J. Mauro. *Solaris Internals: Solaris 10 and OpenSolaris Kernel Architecture*. Prentice Hall, 2nd edition, 2007.

[22] A. Merkel and F. Bellosa. Memory-aware Scheduling for Energy Efficiency on Multicore Processors. In *Proc. of the USENIX Workshop on Power-aware Computing and Systems (HotPower'08)*.

[23] K. K. Rangan, G.-Y. Wei, and D. Brooks. Thread motion: fine-grained power management for multi-core systems. *SIGARCH Comput. Archit. News*, 37(3):302–313, 2009.

[24] S. Rivoire. *Models and Metrics for Energy-efficient Computer Systems*. PhD thesis, Stanford University, 2008.

[25] S. Rivoire, P. Ranganathan, and C. Kozyrakis. A Comparison of High Level Full-System Power Models. In *Proc. of USENIX Workshop on Power Aware Computing and Systems (HotPower'08)*, 2008.

[26] T. Rosing, K. Mihic, and G. De Micheli. Power and reliability management of socs. *Very Large Scale Integration (VLSI) Systems, IEEE Transactions on*, 15(4):391–403, April 2007.

[27] R. Singhal. Inside Intel Next Generation Nehalem Microarchitecture. In *Intel Developer Forum, Shanghai, China*, 2008.

[28] Sun Microsystems. OpenSolaris CPU Power Management - Project Tesla, June 2009.

[29] J. Yang, X. Zhou, M. Chrobak, Y. Zhang, and L. Jin. Dynamic thermal management through task scheduling. *Performance Analysis of Systems and software, 2008. ISPASS 2008. IEEE International Symposium on*, pages 191–201, April 2008.

[30] F. Yao, A. Demers, and S. Shenker. A scheduling model for reduced cpu energy. In *FOCS '95: Proceedings of the 36th Annual Symposium on Foundations of Computer Science (FOCS'95)*, page 374, Washington, DC, USA, 1995. IEEE Computer Society.